

Gan Minghao

☎ (+86) 15218891135 | ✉ liamgan05@gmail.com | 📧 u3667831@connect.hku.hk

EDUCATION

The University of Hong Kong <i>Exchange Student, Department of Engineering</i>	Jan. 2026 – Present Hong Kong
Sun Yat-sen University <i>Bachelor of Engineering in Software Engineering</i> – Academic Records: GPA 4.2/5.0, Rank: 10/231 (Top 5%)	Sep. 2023 – Present Guangzhou, China

SCHOLARSHIP & AWARDS

National Scholarship for Undergraduates , Ministry of Education of China	2023
First Prize, Outstanding Student Scholarship , Sun Yat-sen University	2023
Special Scholarship , School of Software Engineering, Sun Yat-sen University	2023
Second Prize, Outstanding Student Scholarship , Sun Yat-sen University	2024
Fung Scholarship , Fung Foundation & The University of Hong Kong	2025

RESEARCH EXPERIENCE

NAISS Lab, The University of Hong Kong <i>Research Intern (Supervised by Prof. Shinan Liu)</i> – Project: Low Latency Agents System via KV Cache Communication – Designing KV-Cache-communication based agent frameworks to minimize redundant computation. – Based on LMCache, optimizing agent systems through advanced pipeline scheduling and dynamic storage routing. – Engineering KV Cache knowledge fusers and training specialized fusion modules to enable direct, high-fidelity inter-LLM communication via cache states. – Boosting throughput by eliminating redundant text-to-text communication and improving agent performance via the retention of rich latent information in KV caches.	Hong Kong Jan. 2026 – Present
InPlus Lab, Sun Yat-sen University <i>Research Intern (Supervised by Prof. Wuhui Chen)</i> – Project: Low Latency LLM inference system on consumer-grade GPUs via Activation Sparsity – Implemented sparse inference features within the Llama.cpp framework to reduce computation and memory/IO overhead for LLM deployment on consumer GPUs. – Optimized a hybrid computing pipeline enabling asynchronous GPU&CPU execution with on-the-fly weights I/O. – Developed custom CUDA kernels for sparse matrix multiplication, optimizing hardware parallelism. – Demonstrated a 2× speedup in inference throughput over state-of-the-art systems; work submitted to OSDI 2026 .	Guangzhou, China Nov. 2024 – Dec. 2025

SKILLS

Programming : C/C++, CUDA, Python
Frameworks : Llama.cpp, GGML, vLLM, PyTorch
Tools : Git, SSH, Docker, Linux, LaTeX

LANGUAGE PROFICIENCY

IELTS Academic : 7.5 (Listening: 8.0, Reading: 8.5, Writing: 6.5, Speaking: 6.5)